# Ishwar Kalmegh

+91-7218087576 | ishwarkalmegh156@gmail.com | github.com/i5hwar-ka1m39h | linkedin.com/in/ishwar-kalmegh-161004294/

13 October 2026

Dear Hiring Team,

I am applying for the **Cloud Developer / Data Engineer** role with Hewlett Packard Enterprise. With 3+ years of experience building production data systems, the work described here — scalable data pipelines, automated ingestion and processing workflows, and cloud-native solutions that power analytics and AI — is exactly what I do. I am confident working onsite from an HPE office and keen to grow within HPE's edge-to-cloud data platform.

At Techsurvi Pvt. Ltd. I design and build automated data pipelines end to end. Python producers and consumers ingest events onto **Apache Kafka**, **Apache Spark** jobs handle the heavier transformation work, and the results land in **ClickHouse**, **PostgreSQL**, and **S3** on **AWS**. Alongside that I built the **CI/CD** pipelines that unit-test, schedule, and deploy those ETL and application jobs, containerising everything with **Docker** and running it on **Kubernetes**. Before Techsurvi, at Skyfunnel.ai, I built production data and backend services on **PostgreSQL** with **Python** and Node.js, and resolved more than a hundred production data and deployment incidents, cutting resolution time by 40%.

A large part of my work is making data trustworthy and fast. I apply schema and row-level validation, automated backfills, retry/backoff logic, and dead-letter handling on failed consumers so a bad message never silently disappears. I make the platform measurable too: centralised logging through **ELK** and **Grafana**/**Prometheus** dashboards tracking throughput, consumer lag, latency, and error rates, with alerting on each. I also build concurrent **Go** services for load testing and high-volume record validation, and have hands-on experience with **Generative AI** and **Agentic AI** integrations over the same data.

- **Data Engineering:** Kafka, Spark, ClickHouse, ArangoDB, Airflow, ETL/ELT, data quality & validation, optimisation
- **Languages:** Python, Java, Scala, SQL, Bash/shell, Go
- **Cloud & Containers:** AWS (EC2, S3, RDS, Lambda, Glue, Redshift), Azure, GCP, Docker, Kubernetes
- **Databases & Search:** PostgreSQL, MongoDB, Redis, ClickHouse, ArangoDB, DynamoDB, Elasticsearch / OpenSearch
- **Platform:** Linux, shell scripting, CI/CD (Jenkins, GitHub Actions), distributed systems, Grafana, Superset

What appeals to me about HPE is the scope: a global edge-to-cloud platform where the data layer has to work from the edge all the way to the cloud, serving analytics, data science, and next-generation AI applications. I would like to bring that experience to the team, keep deepening my depth in the platform, and contribute to solutions that turn insights into outcomes.

Thank you for your time and consideration. I would welcome the chance to discuss the role.

**Ishwar Kalmegh**
+91-7218087576 | ishwarkalmegh156@gmail.com