Ishwar Kalmegh

Site Reliability Engineer

+91-7218087576 | ishwarkalmegh156@gmail.com | github.com/i5hwar-ka1m39h | linkedin.com/in/ishwar-kalmegh-161004294/

13 October 2026

Dear Mastercard Recruiting Team,

I am applying for the **Site Reliability Engineer II** position on the Enterprise Monitoring team. The dual focus in that description — strategic projects alongside real-time operational needs — is exactly how monitoring work actually lands, and it is how I have worked for the last 3+ years.

I own the production monitoring stack at Techsurvi Pvt. Ltd. **Prometheus**, **Grafana**, **Loki**, **Node Exporter**, and **cAdvisor** — centralising metrics, logs, and dashboards across every service, with alerting on the three things that predict an outage: latency, error rate, and saturation. Alongside running it, I do the other half of the job: day-to-day usage and configuration support for the internal engineering teams who depend on it. That means handling dashboard requests, tuning alerts that are too noisy, and onboarding a new service into the stack properly rather than bolting on an agent and hoping.

I work across the tooling this role names — **Dynatrace**, **AppDynamics**, **Datadog**, and **New Relic** — with **Dynatrace** as my deepest, covering APM instrumentation, custom metrics and services, service maps, distributed tracing, and synthetic monitoring of end-user availability and performance. I am equally comfortable on **Linux** in a terminal, writing and revising automation in **Python**, **Perl**, and **bash**, and running containerized services on **Kubernetes** across **AWS** and **Azure**. Monitoring a container platform means understanding the platform, not just scraping it — agent deployment, probe configuration, and the cluster-versus-workload metric split are all part of the job.

Two things I would bring to an enterprise payments environment specifically. First, discipline under pressure: I have run an **on-call** rotation, and at Skyfunnel.ai I resolved more than a hundred production incidents and cut resolution time by 40% through root cause analysis rather than guesswork. Second, documentation as a deliverable. A monitoring platform is only as good as the team's ability to use it, so I write the technical documentation and runbooks the next person — often me, at 3am — depends on.

I take the Corporate Security Responsibility section seriously as a working constraint, not a formality: monitoring data is among the most sensitive things in an organisation, and least-privilege access, confidentiality of what gets surfaced, and honest reporting of anything anomalous are all part of doing this job well.

- **APM & DEM:** Dynatrace, AppDynamics, Datadog, New Relic, synthetic monitoring, RUM, service maps, distributed tracing
- **Observability:** Prometheus, Grafana, Loki, OpenTelemetry, ELK/OpenSearch, alerting, SLOs and error budgets
- **Platform:** Linux (RHEL, Ubuntu), Kubernetes, Docker, Amazon Web Services, Microsoft Azure
- **Scripting:** Python, Perl, bash, PowerShell; REST API automation, schedulers
- **SRE practice:** on-call rotation, incident response, root cause analysis, runbooks, technical documentation, internal customer support

I'd be glad to talk in more detail about the monitoring platform I've run, and the tooling decisions behind it.

Thank you for your time and consideration.

Ishwar Kalmegh
+91-7218087576 | ishwarkalmegh156@gmail.com