Ishwar Kalmegh

+91-7218087576 | ishwarkalmegh156@gmail.com | github.com/i5hwar-ka1m39h | linkedin.com/in/ishwar-kalmegh-161004294/

13 October 2026

Dear Hiring Team,

I am applying for the **AI-Native SDE-II Engineer** role. With 3+ years of experience building and operating high-scale backend systems, I work in both **Go** and **Node.js/TypeScript** and I own services end to end — design, build, deploy, run, and keep them healthy in production. The two things this role emphasises most, scale and accountability, are the two things I have spent my career building.

At Techsurvi Pvt. Ltd. I design and build high-throughput backend services in Node.js/TypeScript and Go on AWS, holding average API latency at 150 ms. Alongside that I built our **Kafka** event pipelines with partition-keyed ordering, consumer groups, idempotent processing, retry/backoff, dead-letter replay and deliberate backpressure control — the exact failure surface this role calls out. I own the data modelling across **MongoDB**, **PostgreSQL/RDS**, **Redis** and **ClickHouse**, and everything ships through automated CI/CD on Docker. Production is operated, not assumed: centralised metrics, logs and traces in Prometheus, Grafana and Loki, with every change traced back to a dashboard or a trace, not a guess.

Before that, at Skyfunnel.ai, I built the messaging and automation services behind high-volume **WhatsApp** and **Email** delivery — queue-backed processing, scheduling, exponential backoff retries and idempotent handling of redeliveries. I also shipped an **LLM**-powered real-time voice agent end to end: streaming conversations, structured tool calls, and the latency and failure-mode handling that a live conversation demands. That project is how I think about AI work: the model call is the easy part, and the engineering is in the streaming, the fallbacks, the guardrails, and the verification.

On AI-native engineering, I use AI tools for implementation, test generation, debugging and documentation — and I verify the output myself rather than trusting it. Generated code that I cannot explain does not ship. Where it matters, I gate changes with an offline evaluation harness so a prompt or schema change cannot silently degrade behaviour.

- **Core:** Go, Node.js/TypeScript, Kafka/Redpanda, ClickHouse, MongoDB, PostgreSQL/RDS, Redis, AWS
- **Distributed systems:** partitioning, consumer groups, ordering, idempotency, replay, backpressure, queues, schedulers, rate limiting, failure recovery
- **AI/LLM:** RAG, agent workflows, tool/function calling, structured outputs, MCP, evaluation harnesses
- **Production quality:** OpenTelemetry, distributed tracing, metrics/logs/traces, Docker, Kubernetes, CI/CD, incident RCA

I would be glad to walk through the AI product I have built, the architecture behind it, how I evaluate its output, and the trade-offs I made around latency, cost and reliability. Thank you for your time — I would welcome the chance to discuss the role.

Ishwar Kalmegh
+91-7218087576 | ishwarkalmegh156@gmail.com