Ishwar Kalmegh

+91-7218087576 | ishwarkalmegh156@gmail.com | github.com/i5hwar-ka1m39h | linkedin.com/in/ishwar-kalmegh-161004294/

13 October 2026

Dear UBS Recruiting Team,

I am applying for the **Site Reliability Engineer - MQ** role (Job Reference #344023BR) in Pune. I have 3+ years of experience building and operating production hosting infrastructure — the same work this role describes: keeping a messaging platform stable, performant, secure and compliant, and building the hosting and deployment services around it.

At Techsurvi Pvt. Ltd. I build and operate application hosting services using **Docker**, **Docker Swarm** and **Kubernetes**, deploying containerized services with health checks, autoscaling and zero-downtime releases. I own the **GitLab CI/CD** and GitHub Actions pipelines that build, package, test and deploy product artifacts from source through to a running container, with automated test gates and rollback. On the messaging side I support queue-based ingestion paths, automating queue health checks and dead-letter recovery with **Bash** and **Python**, and hardening **Linux** hosts, network security and data encryption to meet policy. I built observability into the platform from the start — centralised metrics, logs and dashboards in **Prometheus**, **Grafana** and **Loki** — so production health is measurable and incidents surface before customers feel them.

Before Techsurvi, at Skyfunnel.ai, I operated the production **message-queue** and messaging services behind **WhatsApp** and **Email** delivery: queue-backed processing, scheduled jobs, exponential backoff retries and dead-letter handling in a high-volume path. That period shaped my incident discipline — I resolved more than a hundred production incidents, cutting resolution time by 40% through systematic debugging, written runbooks and root-cause analysis.

On the AI point in the job description, I use AI to accelerate build, delivery and troubleshooting workflows, but I treat it as a junior analyst rather than an authority. Generated configuration, scripts or code is only used after I have read it, understood why it is correct, and checked it against our standards; nothing reaches a regulated or production path on an unverified suggestion. I document as I go, because in hosting and middleware the person on call at 2am is rarely the person who wrote the change.

- **Messaging:** IBM MQ, Apache Kafka, RabbitMQ, Redis Streams — queue depth, throughput tuning, DLQ recovery, delivery semantics
- **Cloud & Containers:** Microsoft Azure (VMs, AKS, App Service, Functions), AWS, Docker, Kubernetes, Helm
- **Delivery:** GitLab CI/CD, Jenkins, GitHub Actions, Maven/Gradle packaging, artifact repositories, automated testing
- **Observability:** Prometheus, Grafana, Loki, OpenTelemetry, ELK/OpenSearch, incident response and RCA
- **Automation:** Bash, Python, Ansible, Terraform, ARM/Bicep, configuration management

I am genuinely interested in the IBM MQ product and in helping design the next generation cloud-based ecosystem for it — on-premise and in the cloud, built to be easy to consume and efficient to operate. I would welcome the chance to discuss the role and the Pune Middleware crew.

Thank you for your time and consideration.

Ishwar Kalmegh
+91-7218087576 | ishwarkalmegh156@gmail.com